

Microservices Architecture on Energy4Life Monolithic Application

Sunday 9th May, 2021 - 19:40

Kevin L. Biewesch
University of Luxembourg
Email: kevin.biewesch.001@student.uni.lu

Alfredo Capozucca
University of Luxembourg
Email: alfredo.capozucca@uni.lu

Abstract

This document is a template for the scientific and technical (S&T for short) report that is to be delivered by any BiCS student at the end of each Bachelor Semester Project (BSP). The LaTeX source files are available at: <https://github.com/nicolasguelfi/lu.uni.course.bics.global>

This template is to be used using the LaTeX document preparation system or using any document preparation system. The whole document should be in between 6000 to 8000 words¹ (excluding the annexes) and the proportions must be preserved. The other documents to be delivered (summaries, ...) should have their format adapted from this template.

1. Introduction ($\pm 5\%$ of total words)

This paper presents the bachelor semester project made by Motivated Student together with Motivated Tutor as his motivated tutor. It presents the scientific and technical dimensions of the work done. All the words written here have been newly created by the authors and if some sequence of words or any graphic information created by others are included then it is explicitly indicated the original reference to the work reused.

This report separates explicitly the scientific work from the technical one. In deed each BSP must cover those two dimensions with a constrained balance (cf. [BiCS(2018b)]). Thus it is up to the Motivated Tutor and Motivated Student to ensure that the deliverables belonging to each dimension are clearly stated. As an example, a project whose title would be “A multi-user game for multi-touch devices” could define as scientific [Armstrong and Green(2017)] deliverables the following ones:

- Study of concurrency models and their implementation
- Study of ergonomics in human-computer interaction

The length of the report should be from 6000 to 8000 words excluding images and annexes. The sections presenting the technical and scientific deliverables represent $\pm 80\%$ of total words of the report.

1. i.e. approximately 12 to 16 pages double columns

2. Project description ($\pm 10\%$ of total words)

2.1. Domains

2.1.1. Scientific . Provide a description of the scientific domain(s) in which the project is being made.

2.1.2. Technical. Provide a description of the technical domain(s) in which the project is being made.

2.2. Targeted Deliverables

2.2.1. Scientific deliverables. Provide a synthetic and abstract description of the scientific deliverables that were targeted to be produced. Each BSP must contain some work done according to the principles of the scientific method. It basically means that you should define at least one question related to the knowledge domain of your BSP and follow part of the scientific method process to answer to this question. The description of the work done to answer this question is a scientific deliverable.

Examples of question could be:

- Is Python an adequate language for concurrent programs?
- How can we measure the ergonomic of a graphical user interface?
- How can we ensure that a program will not fail?

An answer to such question should be the result of applying partly or totally the scientific method according to its standard definition which can be found in the literature.

As you can see in this template, the scientific deliverable is entirely separated from the technical deliverable. Of course it addresses a question more or less closely related to the technical deliverable.

2.2.2. Technical deliverables. Provide a synthetic and abstract description of the technical deliverables that were targeted to be produced.

3. Pre-requisites ([5%..10%] of total words)

Describe in these sections the main scientific and technical knowledge that is required to be known by you before starting the project. Do not describe in details this knowledge but only abstractly. All the content of this section shall not used, even partly, in the deliverable sections.

3.1. Scientific pre-requisites

3.2. Technical pre-requisites

4. Scientific Deliverable 1 – What is a Microservice?

For each scientific deliverable targeted in section 2 provide a full section with all the subsections described below.

4.1. Requirements ($\pm 15\%$ of section's words)

Since Microservices (MSs) were a previously unknown concept to me, this scientific deliverable targets the familiarisation of this architectural style. In the design section we are going to elaborate on motivations behind the architecture. In the production we will explore what constitutes a MS and the assessment will highlight a few drawbacks of the MS architecture.

4.2. Design ($\pm 30\%$ of section's words)

To understand the benefits and motivations behind using MS over the monolithic style, we will first look at the latter. [1]

Monolithic architectures build applications using a single unit. This basically means that all of an application's functionality, logic, classes, function and name spaces are in the end located within a single executable and thus run within a single process. This model is very much usable in fact and allows us to create functioning systems—as has demonstrated the time of software development before MSs were introduced. However, this architectural style raises some issues nonetheless. For instance, even the smallest of changes in an application require the whole product to be rebuilt and redeployed or redistributed which is not necessarily an ideal scenario.

As the application's life cycles goes on, it is unavoidable to modify it over time. In a monolithic system however, it becomes increasingly difficult to maintain a modular structure—if one was present to begin with—which in turn also leads to modifications becoming more expensive. The expensive modifications are explained by the fact that if a modular structure cannot be guaranteed, it possibly entails that changes to one component may impact other components which were not supposed to be affected. This may result in many adjacent changes and unexpected issues to be performed and treated respectively.

Lastly, if we want to scale an application we actually need to scale the application as a whole rather than the individual components that need to be scaled.

Figure 1 on the following page nicely demonstrates some of the mentioned issues and how the MS architecture tries to tackle them.

Since the application is now split into separate services, making modifications will not require us to build and deploy the whole application anew. It will be enough to only update the services in question.

Further, the fact that the application is built in a modular fashion from the ground up and divided into smaller chunks and services, makes it easier to maintain a modular structure within the application as a whole. As a result, making changes to one such service will guarantee that none of the other services will be affected and thus keep undesired or unforeseen side effects to a minimum.

We can also see that scaling an application becomes much less of a burden. No longer do we need to scale the whole application, but we can simply introduce new services where needed.

4.3. Production ($\pm 40\%$ of section's words)

A first simple definition can be found in the MSs article [1] that lays out some aspects which characterize a MS architecture.

In short, the microservice architectural style is an approach to developing a single application as a suite of small services, each running in its own process and communicating with lightweight mechanisms [...]. These services are built around business capabilities and independently deployable by fully automated deployment machinery. There is a bare minimum of centralized management of these services, which may be written in different programming languages and use different data storage technologies.

However, as stated multiple times in the article [1], there is no formal definition for MS. As such, we shall inspect characteristics that such architectures tend to have in common.

Before we move on, an interesting point to note here is the decreased need for *centralized management*. To give an example [2], centralized management is a structure where only a few individuals make most of the decisions in a company. We can find hierarchies within such a structure, each having to respond to the superior's communicate.

Decentralized management on the other hand would for example allow a manager at a call center or retail store to make instant decisions that impact their work environment. Thus, the decisions are not taken solely by the higher ups but responsibilities are spread across sectors—which could further spread responsibilities. This is an interesting analogy to keep in mind for the following.




Fig. 1. Monoliths and Microservices

4.3.1. Componentization via Services. The article [1] distinguishes between two types of components—where *component* refers to a unit of software that is independently replaceable and upgradeable:

libraries which are being linked into a program and are called using in-memory function calls; and **services** which are external components where communications happen with mechanisms such as web service requests or remote procedure calls.

The advantage of *services* is that such components can be deployed independently—which is not the case with *libraries* because then a component is essentially directly integrated into another component. It should be noted that some service component modifications may affect its interface, which in turn will require other service components to be adjusted accordingly. The aim however is to minimize these through cohesive service boundaries and evolution mechanisms in the service contracts.

Another advantage of services is a more explicit component interface. Often it's only documentation and discipline that prevents clients from breaking a component's encapsulation, thus leading to overly-tight coupling between components. Services make it easier to avoid this by using explicit remote call mechanisms.

4.3.2. Organized around Business Capabilities. Traditionally, splitting a large application into parts is done by assigning a specialized team to each layer of the application, as illustrated in figure 2. This distribution of work force eventually leads to each team worrying only about their specific task instead of the application as a whole, which is an example of *Conway's Law*:

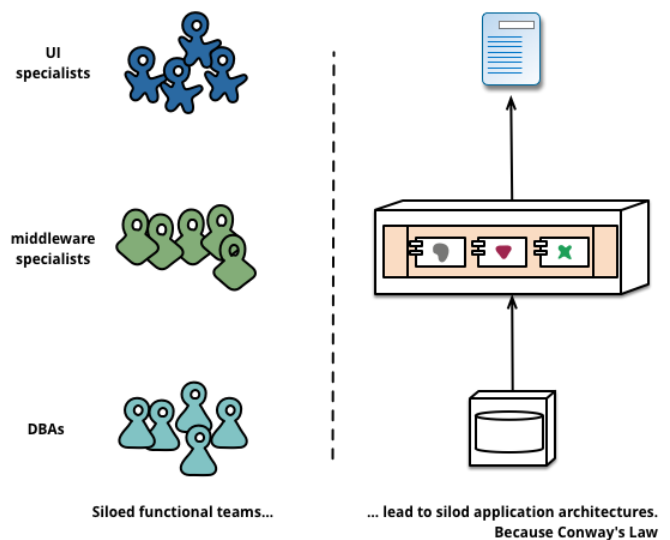


Fig. 2. Conway's Law in action

Any organization that designs a system (defined broadly) will produce a design whose structure is a copy of the organization's communication structure.

In regards to MS, we divide the work force according to *business capabilities* rather than application layers. This implies that our resulting team will be cross-functional², combining the full range of skills required to develop the service in question. This division leads us to figure 3 on the next page.

2. meaning that our team is comprised of people with different expertises, all working closely together for a common goal

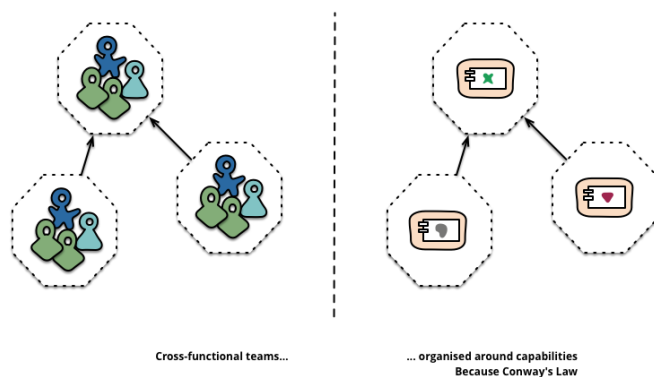


Fig. 3. Service boundaries reinforced by team boundaries

4.3.3. Products not Projects. Most of the application development focuses on delivering a piece of software which is then considered completed. It is then handed over to a separate maintenance team and the original team that built it will then completely abandon the project.

The MS philosophy tries to avoid this approach and instead considers that a team should be responsible for its own product over the full lifetime. Here is also a link to *business capabilities* in that software is no longer considered as a set of functionalities that need to be completed. Rather it is seen as an on-going relationship where one strives to enhance the business capabilities by having the software assist its users.

4.3.4. Smart endpoints and dumb pipes. Often times when it comes to process communication, great importance is put on the communication mechanism itself. With regards to microservices however, we place greater importance on the endpoints rather than the pipes.

Thus one should strive to decouple the MSs and keep them as cohesive as possible. One could almost say that they follow the Unix philosophy in that they receive a request, treat it appropriately, and then return a response. MS can achieve this by using simple REST-like protocols.

This further highlights that services in a microservice architecture are to be isolated components. When they communicate, the magic happens on the services' endpoints and not in the message bus—in fact we need nothing more for messages than simply being transferred from one endpoint to the other.

4.3.5. Decentralized Governance. Centralized governance that often comes along with monolithic architectures have the tendency to standardize on single technology platforms, which can be restricting. Rarely can a single language's full potential be leveraged for the entirety of an application.

Splitting our application into services however, we do have the ability and choice of building each service with whatever tool is best for the job.

This approach also brings a different mindset into development. Instead of focusing solely on defined standards for a product, developers prefer the idea of creating useful tools




Fig. 4. Monolith and Microservice Databases

that can potentially also be used by other developers to solve similar problems.

4.3.6. Decentralized Data Management. Monolithic applications usually prefer to have a single logical database for persistent data, which in turn are then often used across a range of applications. Microservices on the other hand let each service manage its own database. Both situations are illustrated in figure 4.

The issue that arises with decentralizing the data, however, is consistency. Managing this is in fact not an easy task and usually comes at the cost of increased computing times, leading to *eventual consistency* and *compensation operations* being the preferred method. In fact, businesses are often capable of handling a degree of inconsistency in favor of quick responses to demands.

4.3.7. Infrastructure Automation. Many products following the MS architecture make heavy use of Continuous Integration and Continuous Deployment (CI/CD) and infrastructure automation tools. This becomes especially relevant in production when you have to manage a lot of different MSs at once, which simply becomes infeasible without the automation.

4.3.8. Design for failure. One needs to take into consideration that a service may fail at any point in time for whatever reason, and the client should respond to this as gracefully as possible. Due to the decoupled nature of MSs, we have an additional complexity to take into account which constantly requires us to reflect on how service failure can affect the user experience.

It is therefore of utmost importance to detect failures quickly and restore the service. To achieve this, a lot of sophisticated monitoring and logging on various parts of a service are performed.

4.4. Assessment ($\pm 15\%$ of section's words)

Given all the above, there are drawbacks that come with the microservice architecture. As already pointed out previously in section **Design for failure**, a service can fail at any point in

time which makes it quite difficult to build distributed systems. Added to this, remote call are slow which can also have an impact of the final product's performance. Another drawback that we discussed previously is the issue of eventual consistency as mentioned in section **Decentralized Data Management**.

As slightly touched upon in the **Infrastructure Automation** section, we will run into the issue of *operational complexity*³, demanding a mature operations team to manage lots of services, which are being redeployed regularly. A product consisting of half-a-dozen applications can quickly expand into hundreds of little microservices. [3]

Here, the role of DevOps comes into play to ease this whole process. It becomes increasingly harder, tedious and eventually impossible to manage an increasing number of services which also calls for a lot of automation and monitoring to take place. Given that MSs are smaller and thus easier to understand, problems are likely to occur in the interconnections of such services which potentially makes it very difficult to trace and debug issues—hence the need for proper and elaborate monitoring.

[KB:Highlight elements pointing to TD?]

5. Scientific Deliverable 2 – What is the relationship between DevOps and Microservices?

5.1. Requirements

This scientific deliverable aims to find a relationship between the DevOps approach and the MS style. In the design section we are therefore first going to introduce aspects of both DevOps and MSs. The production section will then draw parallels between the presented characteristics which in turn allows us to observe how both concepts are connected to each other. In the assessment we will talk about threats to the validity of our discussions.

5.2. Design

5.2.1. Relevant aspects of microservices. In section 4 on page 2 we have introduced characteristics that generally make up a MS. We shall briefly touch upon some of the ideas again and give a bit more insight that is relevant for our comparison with DevOps.

5.2.1.1. Rapid provisioning. To be able to rapidly launch a server—within a few hours for instance—it is important to automate the process of provisioning⁴. To make serious use MSs, one should strive for a fully automated process—as far as this is possible.

5.2.1.2. Basic monitoring. Due to loosely-coupled services working together in production, things are bound to go wrong in ways that are difficult to detect in test environments. It is therefore of utmost importance to make sure serious

3. while small independent services are easy to deploy, the number of services adds strain to the complexity regarding operations

4. this is the process of adding and configuring tools to a virtual machine

problems—appearing through technical as well as business issues—are detected swiftly.

Technical issues could encompass run-time errors, or service availability. *Business issues* can be related to a drop in orders, or problems with sending transactions through multiple services.

5.2.1.3. Rapid application deployment. A pipeline comprises the building, testing and deploying of an application which are steps that tend to take a long time to run through. The idea is to break up these steps and the tasks within them into stages. While it will add extra time to run through, we benefit from greater confidence because early stages can detect problems that may have devastating effects later down the road. Later stages, in turn, can make more in-depth testing to gain extra confidence in the final product. Stages in a pipeline can be automated or require manual intervention and may even be parallelized—if applicable—to speed up the process. Deployment pipelines are in fact a central part of CI/CD. [4]

Deployment usually involves a pipeline with various steps, each providing some sort of validation of the product. With many services to manage, this pipeline should strive to be fully automated. Achieving this implies close collaboration between developers and operations: the *DevOps Culture* [5].

5.2.2. Relevant aspects of DevOps. Let us now have a look at the DevOps culture before drawing out the relations that exist between the MS and DevOps philosophies. We will again have a look at aspects that are relevant to us but nonetheless do not obfuscate the overall approach that DevOps seeks to follow.

5.2.2.1. The DevOps approach. Traditionally software development is broken down into distinct categories such as requirements analysis, testing and development as well as deployment, operations and maintenance once the product is released. DevOps aims at merging these categories and favour collaboration between **development** and **operations**.

5.2.2.2. Shared responsibilities. A side effect of the DevOps approach are the *shared responsibilities*, encouraging closer collaboration. As briefly mentioned in the **Organized around Business Capabilities** section, now that developers and operators are working under the same hood, the former is much more inclined to adapt a mindset of finding ways to simplify deployment and maintenance since they are closely involved with the tasks of the operators. If the development team were to just hand over the finished product and not worry about it anymore, they would not see the difficulties operators might need to face in light of their productions.

5.2.2.3. Quality of productions. Given that each party is more involved in each other's tasks, this leads to better understanding of the processes and issues which developers and operators need to tackle. This allows for much more efficient problem solving and leads to teams needing to value building quality into the development process. The aim being to automate the deployments and speeding up the testing cycle which results in greater ease of putting code into production.

5.2.2.4. Automation. Automation is an important factor in all of the above. For one, automated tasks—such as testing, configuration and deployment—free staff of these burdens and reduce human error. Another point is that the automation scripts essentially serve as useful and up-to-date documentation of the system. Should, for example, a developer or operator want to inspect or change how a server is configured, they know where to look.

5.3. Production

Given the above points and a few of the characteristics of MSs, one could derive the implication that DevOps is an inherent aspect that comes along with the microservice philosophy. The two points that lead us to this assumption were the **Organized around Business Capabilities** and **Infrastructure Automation** characteristics of MSs.

The infrastructure automation becomes increasingly relevant if our network of MSs expands. Managing only a handful of services by hand can still be feasible, but if we have a few dozens or more services this quickly becomes a virtually impossible task as we would need to take care of each of the countless services individually. As for the DevOps side of things, in order to increase production speed the automation is an equally important factor. Not only does it free the staff of work, but it also reduces human error—which seems likely to seep in when managing lots of services in parallel. These viewpoints on automation allow us to draw an assumption on how the DevOps approach is vital in order to make working with microservices a feasible task.

The second point is that the automation scripts serve as useful and up-to-date documentation of the system. This is very useful in that both developers and operators can inspect these scripts and be on the same page with regards to how the system is built and works. This directly leads us to the MS characteristic of *organizing around business capabilities*. The fact that we do not split our teams according to application layers—as is usually the case for monoliths—but rather according to functionalities greatly favours collaboration of each application layer's expert. The automation scripts coming forth through the DevOps aspect serve in this regard as a bridge to ease this collaboration among team members and thus help further boost and encourage the idea of organising in terms of functionalities rather than isolated application layers.

5.4. Assessment

Possible threats to validity of the above discussion and assumptions stem from the fact that they were drawn with disregard to the technologies. Merely the researched material presented until this point were taken into consideration. Further, we did not have the opportunity to experiment with this organisation of teams in practice, which would surely provide a much better insight and grounds for our assumptions. That said, given more expertise on how MSs and DevOps work in practice might give a more precise and elaborate answer.
[KB:Highlight elements pointing to TD]

6. Technical Deliverable 1 – Does E4L allow for easy deployment of adjacent Microservices?

6.1. Requirements

Our case study consisted in creating a MS relying on the Energy4Life (E4L) application. The latter is built using a monolithic architecture, and we want to find out how easy it is to deploy an adjacent MS application.

The MS to be created should allow us to test the hypothesis that MSs are easy to create and deploy. Further, it should fit into our E4L case study—we will have our service display a graph based on data collected by the E4L application. To close everything off, we shall also create a GitLab pipeline that automates the deployment of this MS.

Hence, the design section will tackle all the steps we had to undertake in order to reach the final and fully automated deployment pipeline. The production section will dive into the technical details and difficulties of each step. The assessment is going to lay out our judgement on how easy it was to deploy the MS given the current E4L application architecture.

6.2. Design

6.2.1. Set up repository. The first step was to set up a repository in which we could place our productions. I had to request an account in order to access the <https://gitlab.uni.lu> GitLab service. An interesting thing to note at this point being that this GitLab service will host our MS code on a different server from the one that contains the E4L code. Thus we can further verify the independence of such services.

6.2.2. Familiarize with docker. We then had to get familiar with the *docker* utility and how to create and manage containers with it. This tool is what would effectively allow us to create isolated services by containerising our application into a docker container.

The rationale behind using docker instead of simply running a process comes back to the points mentioned in the **Composition via Services** section. Simple processes would not enforce the MS architecture as much as containers would, thus allowing for development to shift into a monolithic architecture.

At a later stage we also inspected the *docker-compose* utility which allows for finer and more automated management of containers.

6.2.3. Familiarize with CI/CD. In order to automate the deployment of our application, we want to take advantage of GitLab's CI/CD framework. At this point in time our MS was not fully functioning yet and we were simply experimenting with dummy code. The main goal here was to understand how it works and what it does.

Creating the actual pipeline for our finished MS was only the very last step of the process. The rationale is that automating such a procedure can only be done if one is familiar

with the steps that would otherwise need to be done manually. Therefore it was important to work out everything by hand before finalizing with an automated pipeline.

6.2.4. Obtain data from E4L. Since our MS would display a graph using data from E4L responses, we had to find out how to extract this information. Finding this out was quite an extensive process in itself since the hints were not too clearly laid out in the documentations. More details will be provided in the production.

To give an outline, our first attempt was to sort of brute-force our way into obtain data. While not the most useful of data, it was something to work with. Only later have we discovered that E4L exposes an API which greatly alleviates the burden of our brute-force approach.

Obtaining useful data however, was only achievable thanks to the help and guidance of a prior team member of E4L who had great responsibility in creating the back-end. He showed me the proper way to obtain all the data from the API as well as by directly connecting to the E4L database.

6.2.5. Write code which retrieves data and displays it.

For this purpose we relied on productions made during my summer job where I was responsible for finding ways to plot graphs. For the sake of simplicity we used Python to write this code. It should be noted that this and the previous step of obtaining data were always running in parallel.

In the end, it was simply a matter of retrieving the data from the API or the database, and then to plot some data using example code produced during the summer job. It should further be noted that we were not deeply interested in *what* data is being displayed, it should rather serve as a proof of concept for our purpose.

6.2.6. Containerize and deploy the application. Once the application was working locally, we attempted to make it run inside a docker container. Note that only the version retrieving data from the API was used for experimenting this way. The database version is analogous and did not require great modifications. Further it had a restriction where the database was only accessible to certain docker containers running on the server where the database lives. We will provide more details in the production.

Containerizing the application came with a fair share of issues however both locally as well as once deployed on the testing server. Work-arounds and fixes were implemented and resulted in a fully functional service.

6.2.7. Create pipeline to automate deployment. The final step was to create a pipeline that automates the whole process. Since we were making a Python application, there were no building stages that had to be made in the pipeline. Thus we only ended up with a deployment stage which was very short thanks to the docker-compose utility taking care of everything. One could possibly add a testing stage to verify availability of

E4L—since our MS cannot run without being able to probe E4L for data—but this was not implemented for this project. **[KB:Methodlogy: Steps required to achieve final solution]**

6.3. Production

[KB:More details on each step + difficulties encountered]

6.4. Assessment

[KB:Was is easy to deploy MS? Flaws in the service that still need to be treated? (i.e. availability of E4L)]

Acknowledgment

The authors would like to thank the BiCS management and education team for the amazing work done.

7. Conclusion

The conclusion goes here.

References

- [BiCS(2018a)] BiCS Bachelor Semester Project Report Template. <https://github.com/nicolasguelfi/lu.uni.course.bics.global> University of Luxembourg, BiCS - Bachelor in Computer Science (2017).
- [BiCS(2018b)] Bachelor in Computer Science: BiCS Semester Projects Reference Document. Technical report, University of Luxembourg (2018)
- [Armstrong and Green(2017)] J Scott Armstrong and Kesten C Green. Guidelines for science: Evidence and checklists. *Scholarly Commons*, pages 1–24, 2017. https://repository.upenn.edu/marketing_papers/181/
- [1] Microservices: A definition of this new architectural term <https://martinfowler.com/articles/microservices.html>
- [2] Centralized and Decentralized Management Explained <https://content.personalfinancelab.com/finance-knowledge/management/centralized-and-decentralized-management-explained/?v=c4782f5abe5c>
- [3] Microservice Trade-Offs – Operational Complexity (Con) <https://martinfowler.com/articles/microservice-trade-offs.html#ops>
- [4] DeploymentPipeline <https://martinfowler.com/bliki/DeploymentPipeline.html>
- [5] DevOpsCulture <https://martinfowler.com/bliki/DevOpsCulture.html>
- [6] [;VIM;] []VIM{}

8. Appendix

All images and additional material go there.